

Anomaly Detection in Graphs

Coen Valk

1 Introduction

Anomaly detection is an important topic in many fields including financial fraud detection, event tracking, and network security to name a few. Several methods for detecting anomalies have been created over the years, each with its advantages and disadvantages. Here we will discuss the performance of certain anomaly detection algorithms on a few different datasets to learn more about the strengths and weaknesses of each algorithm.

2 Background

Anomalies can be found at several different scopes of a graph:

- **Global:** Where each node is compared to every other node, without regard for graph structure.
- **Local:** Where each node is compared to its direct neighbors.
- **Community:** Where each node is compared to nodes in the same community.

Searching for anomalies at any of these three scopes may produce differing results. For instance, if one were to have a graph of interactions between people with node attributes on income, a global outlier detection method would consider such high wealth individuals including Jeff Bezos and Bill Gates as anomalies. However, within the context of the billionaires that Jeff Bezos and Bill Gates typically interact with - by comparing nodes at the local or community scale - their income might not be as anomalous as we had originally thought.

2.1 GLODA

the global outlier detection algorithm (GLODA) finds outliers without considering graph structure. This is often done using the local outlier factor (LOF) algorithm [3]. LOF determines outliers by identifying how many k-nearest neighbors are in common between nodes. if a node has few neighbors in common with other nodes, it will be considered anomalous.




Figure 1: GLODA outlier scores

Figure 1 shows how GLODA determines outliers in three different graph topologies. Each node in the figure is labeled with its attribute (in this case an array containing one value) and colored with its corresponding outlier score. Observe that GLODA does not find any anomalies in the ring of cliques example, as there are several nodes with attribute 100, and without the graph context that might not seem peculiar.

2.2 DNODA

local outliers can be found using the direct neighbor outlier detection algorithm (DNODA). It scores each node based on the average distance from its direct neighbors given by the equation:

$$\text{Score}(v) = \frac{\sum_{u \in N_v} \text{dist}(u, v)}{|N_v|}$$

Where N_v is the set of neighbors for node v .




Figure 2: DNODA outlier scores

Figure 2 shows DNODA's labeling on the three graphs. observe that the star graph is labeled completely anomalous. Because each node is connected to the center, the mean distance of neighbors is the same for each node - so each node is given a distance of 100. This result is not usable, and so all nodes would be considered not anomalous. On the other hand, DNODA does well identifying anomalous nodes in the cycle and ring of cliques structures.

2.3 CNA

The community neighbor algorithm (CNA) is very similar to DNODA except its scope broadens to a community. CNA needs to first perform a community detection algorithm to find clusters of nodes, then for each node finds the average distance between itself and the rest of the nodes in its community as shown in the equation:

$$\text{Score}(v) = \frac{\sum_{u \in C_v} \text{dist}(u, v)}{|C_v|}$$

Where C_v is the community that contains node v .




Figure 3: CNA outlier scores

Figure 3 shows CNA’s labeling for each of the three graph structures. CNA identifies anomalies in these three graph examples rather well.

2.4 Oddball

Another common method for anomaly detection as is used by the Oddball algorithm is coming up with patterns in a dataset including relationships between nodes and edges, edges and total edge weights. Deviations from this pattern could be considered anomalous [2]. Oddball works on the theory that most graphs exhibit power law relationships between certain graph attributes, even at a local neighbor scale. Oddball performs this power law analysis on each node’s egonet - the subgraph containing all of a nodes direct neighbors and edges between such neighbors. Oddball’s outlier score function is

$$\text{Score}(i) = \frac{\max(y_i, Cx_i^\theta)}{\min(y_i, Cx_i^\theta)} \cdot \log(|y_i - Cx_i^\theta| + 1)$$

Where C, θ are the best fit parameters that fit the power law relationship between x, y .

As Oddball works on the assumption that the input graph exhibits power law relationships, the three simple example graph topologies used to demonstrate performance of GLODA, DNODA and CNA are not effective. Instead, we will construct a small graph that does exhibit real-world observed power law relationships with respect to egonet-density power law (EDPL) and edge-weight power law (EWPL) to showcase how Oddball differs from GLODA, DNODA, and CNA [1].

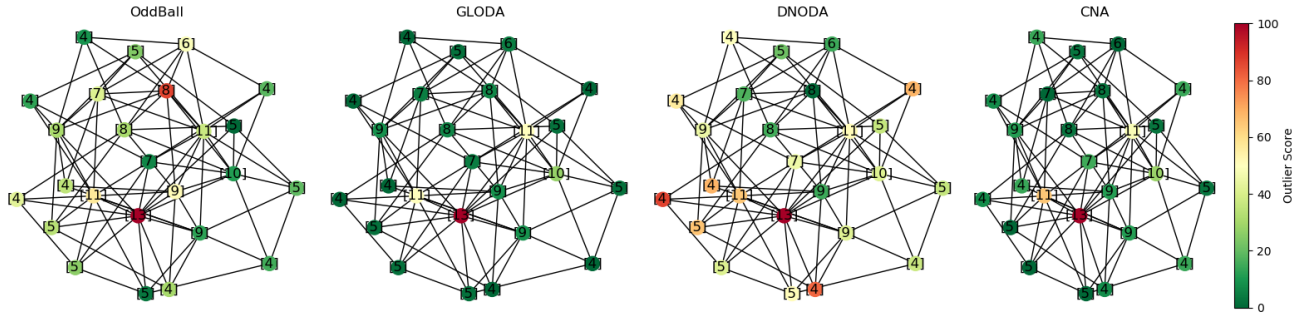


Figure 4: Oddball vs. GLODA, DNODA, and CNA

Observe in figure 4 how certain nodes are colored differently between algorithms. Notably, the Oddball algorithm finds outliers that the other 3 algorithms do not pick up on. These are nodes that fall outside of the typical expectation of the power law in its egonet, but might not specifically be distant from other nodes.

3 Methods

Each algorithm will run on 3 datasets to identify nodes that all algorithms identify as anomalous and nodes that only one algorithm considers anomalous. Then we will combine the results of all of these algorithms’ rankings using the borda voting system to come up with one list of anomalous nodes. Some information on the three datasets:

- **Enron:** Emails between executives at Enron.
- **FEC donation data:** Federal Election Committee data of donations from individuals to political committees.
- **Elliptical Bitcoin Transaction Data:** Financial transaction network with labeled data on ”illicit” nodes. The behavior of illicit nodes is more nuanced than distance from the average.

Each of these algorithms are node-focused and GLODA, DNODA, and CNA depend on distance between nodes to determine outlier score. Therefore each dataset has features that need to be engineered slightly before each algorithm can run properly on the data.

3.1 Enron

The only feature engineering for the Enron dataset for this example is the degree per node. This will find anomalies that have significantly higher or lower degree than the other nodes it is compared to. This will be necessary for GLODA, DNODA and CNA to make inference on anomalies.

To match this outlier detection attempt, the Oddball algorithm will fit to the edge density power law (EDPL) to find either near-cliques or near-stars.

3.2 FEC Donation Data

The FEC donation data has edge weights for donation amount. As mentioned earlier however, these algorithms are node-focused. To account for this, each node will have as attributes the mean and median donation from that node. This will find nodes that have particularly large or particularly small donations compared to similar nodes.

For the Oddball algorithm, we will fit to the edge-weight power law (EWPL) to find similar anomalies in the graph.

As an added note, if it is necessary to do anomaly detection on edges like in credit card fraud detection, one may perform the same anomaly detection algorithms on the induced line graph of the data. Unfortunately with large hubs such as in the donation network, a line graph will have large cliques which are prohibitively expensive to work with. For this example we will stick with data from the original graph structure.

3.3 Bitcoin Transaction Data

The Bitcoin transaction dataset has more than 100 numerical attributes per node. To save on compute time the dataset is reduced using principal component analysis (PCA) while keeping 80% of the variance. This reduces the dimensionality all the way to 13 numeric attributes per node.

4 Results

4.1 Enron

Enron Worst Offenders:

```
GLODA: [4938 3994 6374 8876 2598 280 1 6211 1356 3982 4406 1470 4281 1394
4283 2375 159 992 3052 821] 0:00:02.054511
DNODA: [27137 27167 27148 27150 27164 27163 27162 27161 27160 27159 27149 27158
27156 27155 27154 27153 27152 27151 27157 27646] 0:00:01.233046
CNA: [ 107 97 95 188 244 89 182 93 1273 96 92 79 148 245
85 197 80 191 144 271] 0:00:21.412095
OddBall: [34257 21043 32192 33284 30623 36248 34797 36640 30609 35234 33276 224
22412 36324 271 36311 153 36297 20260 32713] 0:00:14.729157
Borda: [244 96 148 153 119 197 80 191 144 271]
```

Done.

Above are the top 20 outliers for each algorithm when run on the Enron dataset along with the result after ranking with borda. Notice that there are a couple similarities between the outputs of each algorithm. It seems that in this example, the community neighbor algorithm performed well by identifying anomalous nodes that were generally agreed upon by all of the algorithms to be outliers.

4.2 FEC Donation Data

FEC Donation Data Worst Offenders:

```
GLODA: [ 55 34519 60281 60737 20017 50180 1675 2558 14665 56009 2060 34646
29594 36236 43165 41717 25511 54672 28561 43911] 0:00:03.711894
DNODA: [30013 29406 60907 6780 28342 31234 42065 15983 10335 57382 23679 67106
67695 29179 48156 67911 14308 9859 46801 4893] 0:00:00.981846
CNA: [28342 10335 7186 54303 29179 48156 67911 57382 30013 6780 31234 15983
42065 23679 67695 67106 4893 46801 14308 9859] 0:05:31.636945
OddBall: [50548 9825 21396 21340 67695 17148 29301 44857 1480 67911 29179 48156
```

```

28342 3574 47968 47862 29406 60907 46801 4893] 0:11:41.377449
Borda: [25664 56528 45919 8450 49473 49891 233 25511 9822 22647]
Done.

```

In this example, DNODA and Oddball most notably have some identified nodes in common, albeit in a slightly different order. Beyond this, there are no node ids that stand out as being in every algorithm’s list of highest outlier score.

4.3 Bitcoin Transaction Data

Bitcoin Data Worst Offenders:

```

GLODA: [132852 17622 4265 133605 12532 11652 11025 129090 134619 7215
6391 128164 127304 4486 3378 5913 7521 1376 1661 136165] 0:00:36.455671
DNODA: [191731 192356 190412 192608 193037 193137 192737 192039 192919 192728
192504 192282 192021 190783 144268 192979 193021 140237 140236 141690] 0:00:03.762974
CNA: [144268 144269 115222 126218 116671 111866 111967 127239 63415 203550
193603 126265 126344 140236 141690 127504 126639 189851 197761 140237] 0:16:15.680851
OddBall: [192039 190412 192356 191596 192015 192919 192282 190783 191731 192737
193021 192728 192608 193137 193037 192979 192504 130048 140236 141690] 0:00:29.681412
Borda: [126344 197761 115008 126639 127504 126517 126265 140237 140236 141690]
Done.

```

This dataset showed the most consistent results, with certain nodes being in almost all of the algorithm’s top 20 outliers. This results in a borda count very similar to both Oddball as well as DNODA.

5 Conclusion

Ultimately, each algorithm does not necessarily perform ”better” or ”worse” than any of the others, but simply looks for different kinds of anomalies, and at times all of these algorithms produce similar results. The use of any of these algorithms is determined mostly by the user’s needs and what kind of anomalous behavior they are looking for.

References

- [1] L. Akoglu, M. McGlohon, and C. Faloutsos. Rtm: Laws and a recursive generator for weighted time-evolving graphs. In *2008 Eighth IEEE International Conference on Data Mining*, pages 701–706, 2008.
- [2] Leman Akoglu, Mary McGlohon, and Christos Faloutsos. Oddball: Spotting anomalies in weighted graphs. pages 410–421, 07 2010.
- [3] Markus Breunig, Hans-Peter Kriegel, Raymond Ng, and Joerg Sander. Lof: Identifying density-based local outliers. volume 29, pages 93–104, 06 2000.